

Building and Evaluating an End-to-end Information Retrieval (IR) System

Computing Lab Examination, CSE, IIT Kharagpur

Full Marks: 75

Time allotted: 2 hours

1. **Using the Vector Space Model (VSM).** In this model, the term frequency and term weight values for each document (and query) can be represented (and stored) as vectors; hence the name vector space model. [5 x 6 = 30]

- Download and *preprocess* (as discussed in class) the IR resources (queries, document collection and relevance judgments) that have been previously provided to you. [5]
- As earlier, build inverted indices for each term in the collection. Now apply the vector space model for document retrieval as follows. For each term-document pair (t, d) , compute the term frequencies $tf_{t,d}$ where $tf_{t,d}$ is the number of times t appears in d . [5]
- For each term t , compute the Inverse Document Frequency idf_t as follows.

$$idf_t = \log_{10} \frac{N}{df_t}$$

where N is the total number of documents in the collection and df_t is the document frequency of t , i.e. the number of documents in which t occurs. [5]

- Compute the tf-idf (also called term weights in this context) for *every* term-document pair as follows (the small hyphen between tf and idf in the tf-idf term is **not** a minus sign). [5]

$$tf-idf_{t,d} = tf_{t,d} \times idf_t$$

- Compute the overlap score for *every* query-document pair as described below. [5]

$$Score(q, d) = \sum_{t \in q} tf-idf_{t,d}$$

- Then, for each query, find and output the ten documents that have the highest overlap score with the query. [5]

2. **Evaluation of a ranked list: Mean Average Precision (MAP).** MAP is one of the popular metrics to evaluate ranked lists with binary relevance levels. [5 + 25 + 5 = 35]

- Convert relevance judgments in the *qrels* files 1 and 2 to 1 (relevant), and 0 as non-relevant. [5]
- Compute Average Precision (AP) for each query q as follows for the ranked result list containing ten documents for each query derived in Question 1f. [25]

$$AP(q) = \frac{\sum_{k=1}^{10} (Precision(k) \times rel(k))}{\#(relevant\ documents\ for\ q)}$$

where k is a position in the ranked list ($1 \leq k \leq 10$), $rel(k)$ is the relevance of the document at rank k with respect to q (i.e. $rel(k)$ is either 0 or 1), and $Precision(k)$ is given by

$$Precision(k) = \frac{\#(\text{relevant documents retrieved till rank } k)}{k}$$

- c. Compute MAP for your result set (i.e. each query accompanied by the top-10 documents) as follows:

$$MAP(Q) = \frac{1}{|Q|} \sum_{j=1}^{|Q|} AP(q_j)$$

where Q is the set of all queries.

- 3. Evaluation of a ranked list: Discounted Cumulative Gain (DCG@10).** DCG (and its normalized variant nDCG) is the most widely used metric today for evaluating ranked result lists with graded relevance levels (0, 1, 2 in this case). For computing DCG, use the original relevance judgments of 0, 1, and 2 (and **not** the binary levels as in Q. 2). For each query q , find the DCG@10 (Discounted Cumulative Gain) for its graded ranked list of 10 documents as follows

$$DCG@10 = rel_1 + \sum_{i=2}^{10} \frac{rel_i}{\log_2(i+1)}$$

where rel_i is the graded relevance judgment (0, 1 or 2) of the document at Rank i in the result list. Compute the average DCG@10 for all the queries in Q . **[10]**